

fc_tune_vgg16_2class

January 21, 2020

```
[1]: import numpy as np
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dropout, Flatten, Dense
from tensorflow.keras import applications
from tensorflow.keras.utils import to_categorical
from tensorflow.keras import optimizers
from tensorflow.keras.utils import plot_model
import matplotlib.pyplot as plt
```

```
[24]: # dimensions of our images.
img_width, img_height = 224, 224

top_model_weights_path = '/home/user/models/top_tuned/
↳bottleneck_fc_model_2class.h5'
train_data_dir = '/home/user/      /convnets/transfer-learning-keras/dataset/
↳training'
validation_data_dir = '/home/user/      /convnets/transfer-learning-keras/
↳dataset/validation'
evaluation_data_dir = '/home/user/      /convnets/transfer-learning-keras/
↳dataset/evaluation'
nb_train_samples = 3000
nb_validation_samples = 1000
nb_evaluation_samples = 1000
epochs = 20
batch_size = 20
```

```
[3]: def save_bottleneck_features():
    datagen = ImageDataGenerator(rescale=1. / 255)

    # build the VGG16 network
    model = applications.VGG16(include_top=False, weights='imagenet')

    generator = datagen.flow_from_directory(
        train_data_dir,
        target_size=(img_width, img_height),
        batch_size=batch_size,
```

```

        class_mode=None,
        shuffle=False)
bottleneck_features_train = model.predict_generator(
    generator, nb_train_samples // batch_size)
np.save(open('bottleneck_features_train.npy', 'wb'),
        bottleneck_features_train)

generator = datagen.flow_from_directory(
    validation_data_dir,
    target_size=(img_width, img_height),
    batch_size=batch_size,
    class_mode=None,
    shuffle=False)
bottleneck_features_validation = model.predict_generator(
    generator, nb_validation_samples // batch_size)
np.save(open('bottleneck_features_validation.npy', 'wb'),
        bottleneck_features_validation)

```

```

[19]: def train_top_model():
    train_data = np.load(open('bottleneck_features_train.npy', 'rb'))
    train_labels = np.array(
        [0] * (nb_train_samples // 2) + [1] * (nb_train_samples // 2))

    validation_data = np.load(open('bottleneck_features_validation.npy', 'rb'))
    validation_labels = np.array(
        [0] * (nb_validation_samples // 2) + [1] * (nb_validation_samples // 2))

    model = Sequential()
    model.add(Flatten(input_shape=train_data.shape[1:]))
    model.add(Dense(256, activation='relu'))
    model.add(Dropout(0.5))
    model.add(Dense(1, activation='sigmoid'))

    model.compile(optimizer='adam',
                  loss='binary_crossentropy', metrics=['accuracy'])

    history = model.fit(train_data, train_labels,
                        epochs=epochs,
                        batch_size=batch_size,
                        validation_data=(validation_data, validation_labels))
    model.save(top_model_weights_path)
    return model, history

```

```

[11]: def plot(model, history):
    plot_model(model, to_file='model.png')
    # Plot training & validation accuracy values
    plt.plot(history.history['accuracy'])

```

```

plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

# Plot training & validation loss values
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```

```
[6]: save_bottleneck_features()
```

Found 3000 images belonging to 2 classes.
Found 1000 images belonging to 2 classes.

```
[20]: model, history = train_top_model()
```

```

Train on 3000 samples, validate on 1000 samples
Epoch 1/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.4256 -
accuracy: 0.8817 - val_loss: 0.1881 - val_accuracy: 0.9180
Epoch 2/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.1437 -
accuracy: 0.9447 - val_loss: 0.1149 - val_accuracy: 0.9510
Epoch 3/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0995 -
accuracy: 0.9613 - val_loss: 0.1120 - val_accuracy: 0.9540
Epoch 4/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0799 -
accuracy: 0.9710 - val_loss: 0.0927 - val_accuracy: 0.9640
Epoch 5/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0464 -
accuracy: 0.9840 - val_loss: 0.1135 - val_accuracy: 0.9600
Epoch 6/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0467 -
accuracy: 0.9810 - val_loss: 0.1467 - val_accuracy: 0.9410
Epoch 7/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0286 -
accuracy: 0.9907 - val_loss: 0.1198 - val_accuracy: 0.9650
Epoch 8/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0230 -



```

```

accuracy: 0.9910 - val_loss: 0.1097 - val_accuracy: 0.9630
Epoch 9/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0217 -
accuracy: 0.9927 - val_loss: 0.1227 - val_accuracy: 0.9600
Epoch 10/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0311 -
accuracy: 0.9873 - val_loss: 0.1240 - val_accuracy: 0.9630
Epoch 11/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0283 -
accuracy: 0.9917 - val_loss: 0.1434 - val_accuracy: 0.9580
Epoch 12/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0186 -
accuracy: 0.9937 - val_loss: 0.1604 - val_accuracy: 0.9620
Epoch 13/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0284 -
accuracy: 0.9883 - val_loss: 0.1567 - val_accuracy: 0.9600
Epoch 14/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0397 -
accuracy: 0.9840 - val_loss: 0.1226 - val_accuracy: 0.9630
Epoch 15/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0174 -
accuracy: 0.9947 - val_loss: 0.1513 - val_accuracy: 0.9610
Epoch 16/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0281 -
accuracy: 0.9890 - val_loss: 0.1444 - val_accuracy: 0.9630
Epoch 17/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0199 -
accuracy: 0.9937 - val_loss: 0.1518 - val_accuracy: 0.9660
Epoch 18/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0395 -
accuracy: 0.9850 - val_loss: 0.1558 - val_accuracy: 0.9600
Epoch 19/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0111 -
accuracy: 0.9967 - val_loss: 0.2160 - val_accuracy: 0.9610
Epoch 20/20
3000/3000 [=====] - 4s 1ms/sample - loss: 0.0253 -
accuracy: 0.9893 - val_loss: 0.1597 - val_accuracy: 0.9550

```

```
[21]: plot(model, history)
```



```
[32]: def evaluation(model):  
    datagen = ImageDataGenerator(rescale=1. / 255)  
    generator = datagen.flow_from_directory(  
        evaluation_data_dir,  
        target_size=(img_width, img_height),  
        batch_size=batch_size,  
        class_mode=None,  
        shuffle=False)  
    model1 = applications.VGG16(include_top=False, weights='imagenet')  
    features = model1.predict_generator(  
        generator, nb_evaluation_samples // batch_size)  
    return features
```

```
[33]: features = evaluation(model)
```

Found 1000 images belonging to 2 classes.

```
[38]: evaluation_labels = np.array(  
    [0] * (nb_evaluation_samples // 2) + [1] * (nb_evaluation_samples // 2))  
model.test_on_batch(features, evaluation_labels)
```

```
[38]: [0.20799212, 0.952]
```