

fc_tune_vgg16_2class

January 21, 2020

```
[1]: import numpy as np
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dropout, Flatten, Dense
from tensorflow.keras import applications
from tensorflow.keras.utils import to_categorical
from tensorflow.keras import optimizers
from tensorflow.keras.utils import plot_model
import matplotlib.pyplot as plt

[2]: # dimensions of our images.
img_width, img_height = 224, 224

top_model_weights_path = '/home/user/models/top_tuned/
↳bottleneck_fc_model_2class.h5'
train_data_dir = '/home/user/      /convnets/transfer-learning-keras/dataset/
↳training'
validation_data_dir = '/home/user/      /convnets/transfer-learning-keras/
↳dataset/validation'
nb_train_samples = 3000
nb_validation_samples = 1000
epochs = 20
batch_size = 20

[3]: def save_bottleneck_features():
    datagen = ImageDataGenerator(rescale=1. / 255)

    # build the VGG16 network
    model = applications.VGG16(include_top=False, weights='imagenet')

    generator = datagen.flow_from_directory(
        train_data_dir,
        target_size=(img_width, img_height),
        batch_size=batch_size,
        class_mode=None,
        shuffle=False)
    bottleneck_features_train = model.predict_generator(
```

```

        generator, nb_train_samples // batch_size)
np.save(open('bottleneck_features_train.npy', 'wb'),
        bottleneck_features_train)

generator = datagen.flow_from_directory(
    validation_data_dir,
    target_size=(img_width, img_height),
    batch_size=batch_size,
    class_mode=None,
    shuffle=False)
bottleneck_features_validation = model.predict_generator(
    generator, nb_validation_samples // batch_size)
np.save(open('bottleneck_features_validation.npy', 'wb'),
        bottleneck_features_validation)

```

```

[4]: def train_top_model():
    train_data = np.load(open('bottleneck_features_train.npy', 'rb'))
    train_labels = np.array(
        [0] * (nb_train_samples // 2) + [1] * (nb_train_samples // 2))

    validation_data = np.load(open('bottleneck_features_validation.npy', 'rb'))
    validation_labels = np.array(
        [0] * (nb_validation_samples // 2) + [1] * (nb_validation_samples // 2))

    model = Sequential()
    model.add(Flatten(input_shape=train_data.shape[1:]))
    model.add(Dense(256, activation='relu'))
    model.add(Dropout(0.5))
    model.add(Dense(1, activation='sigmoid'))

    model.compile(optimizer='rmsprop',
                  loss='binary_crossentropy', metrics=['accuracy'])

    history = model.fit(train_data, train_labels,
                        epochs=epochs,
                        batch_size=batch_size,
                        validation_data=(validation_data, validation_labels))
    model.save(top_model_weights_path)
    return model, history

```

```

[11]: def plot(model, history):
    plot_model(model, to_file='model.png')
    # Plot training & validation accuracy values
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('Model accuracy')
    plt.ylabel('Accuracy')

```

```
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

# Plot training & validation loss values
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```

```
[6]: save_bottleneck_features()
```

Found 3000 images belonging to 2 classes.
Found 1000 images belonging to 2 classes.

```
[7]: model, history = train_top_model()
```



```
Train on 3000 samples, validate on 1000 samples
Epoch 1/20
3000/3000 [=====] - 8s 3ms/sample - loss: 0.8015 -
accuracy: 0.8477 - val_loss: 0.3649 - val_accuracy: 0.8830
Epoch 2/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.2750 -
accuracy: 0.9153 - val_loss: 0.1310 - val_accuracy: 0.9510
Epoch 3/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.1901 -
accuracy: 0.9403 - val_loss: 0.2079 - val_accuracy: 0.9450
Epoch 4/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.1524 -
accuracy: 0.9503 - val_loss: 0.1793 - val_accuracy: 0.9620
Epoch 5/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.1259 -
accuracy: 0.9623 - val_loss: 0.1692 - val_accuracy: 0.9510
Epoch 6/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0967 -
accuracy: 0.9690 - val_loss: 0.1397 - val_accuracy: 0.9600
Epoch 7/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.1002 -
accuracy: 0.9727 - val_loss: 0.2300 - val_accuracy: 0.9580
Epoch 8/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0527 -
accuracy: 0.9820 - val_loss: 0.2549 - val_accuracy: 0.9600
Epoch 9/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0682 -
```

```

accuracy: 0.9833 - val_loss: 0.4271 - val_accuracy: 0.9500
Epoch 10/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0513 -
accuracy: 0.9853 - val_loss: 0.1966 - val_accuracy: 0.9620
Epoch 11/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0367 -
accuracy: 0.9897 - val_loss: 0.2310 - val_accuracy: 0.9620
Epoch 12/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0459 -
accuracy: 0.9870 - val_loss: 0.2495 - val_accuracy: 0.9580
Epoch 13/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0381 -
accuracy: 0.9903 - val_loss: 0.2772 - val_accuracy: 0.9600
Epoch 14/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0401 -
accuracy: 0.9890 - val_loss: 0.2559 - val_accuracy: 0.9640
Epoch 15/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0227 -
accuracy: 0.9913 - val_loss: 0.3340 - val_accuracy: 0.9610
Epoch 16/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0250 -
accuracy: 0.9933 - val_loss: 0.3610 - val_accuracy: 0.9600
Epoch 17/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0222 -
accuracy: 0.9927 - val_loss: 0.5701 - val_accuracy: 0.9470
Epoch 18/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0299 -
accuracy: 0.9943 - val_loss: 0.3645 - val_accuracy: 0.9570
Epoch 19/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0219 -
accuracy: 0.9940 - val_loss: 0.4468 - val_accuracy: 0.9590
Epoch 20/20
3000/3000 [=====] - 7s 2ms/sample - loss: 0.0220 -
accuracy: 0.9947 - val_loss: 0.4269 - val_accuracy: 0.9550

```

```
[12]: plot(model, history)
```



[]: